

Lab1

September 28, 2016

1 Akari Ishikawa RA163282

2 Exercício 01 - MC886 A

```
In [1]: import numpy as np
import pandas as pd
```

```
In [2]: import sklearn
import sklearn.discriminant_analysis as lda
import sklearn.decomposition
```

```
In [3]: data = np.genfromtxt('data.csv', delimiter=',')
print(data.shape)
```

(477, 167)

```
In [4]: # Removemos a última coluna do csv e a primeira linha que falava das features
data1 = data[1:, :-1]
print(data1.shape)
```

(476, 166)

3 1) PCA

```
In [5]: pca = sklearn.decomposition.PCA(copy=True)
```

```
In [6]: pca.fit(data1)
```

```
Out[6]: PCA(copy=True, n_components=None, whiten=False)
```

```
In [7]: np.cumsum(pca.explained_variance_ratio_)
```

```
Out[7]: array([ 0.31188125,  0.45110049,  0.52728075,  0.57871107,  0.62790775,
                0.66868927,  0.70115259,  0.73157891,  0.75149626,  0.76866767,
                0.7841581 ,  0.79847584,  0.81177395,  0.8244771 ,  0.83539864,
                0.84538246,  0.85404078,  0.86220635,  0.8697735 ,  0.87691348,
```

```

0.88366985, 0.89003952, 0.89635954, 0.90219742, 0.90778833,
0.91302009, 0.91796749, 0.92286063, 0.92753044, 0.93196024,
0.93596809, 0.93996068, 0.94380535, 0.94713044, 0.95006551,
0.952904 , 0.95556949, 0.95810425, 0.96053757, 0.96258962,
0.96459123, 0.96645684, 0.96824055, 0.96989498, 0.97151772,
0.97300967, 0.97444944, 0.97580079, 0.97703677, 0.97823005,
0.97936854, 0.98044865, 0.98148826, 0.98244777, 0.98335119,
0.98421892, 0.98504699, 0.98582422, 0.98658195, 0.98729275,
0.98793209, 0.98855017, 0.98909833, 0.98964073, 0.99015928,
0.99062971, 0.99108475, 0.99150553, 0.99190155, 0.9922765 ,
0.99262038, 0.99295524, 0.99328065, 0.99360092, 0.99391974,
0.99421253, 0.994494 , 0.99475744, 0.99501492, 0.99526007,
0.99548223, 0.995696 , 0.99590043, 0.99610138, 0.99628465,
0.99646083, 0.99663284, 0.99680215, 0.99696691, 0.99711137,
0.99724949, 0.99738161, 0.99750783, 0.99762882, 0.9977424 ,
0.99785097, 0.99795661, 0.99805876, 0.99815604, 0.99825005,
0.99834086, 0.99842216, 0.99850204, 0.99858134, 0.99865945,
0.99873451, 0.99880654, 0.99887223, 0.99893224, 0.99899107,
0.99904907, 0.99910407, 0.99915774, 0.99920611, 0.99925271,
0.99929426, 0.99933471, 0.99937279, 0.99940891, 0.9994431 ,
0.99947583, 0.99950773, 0.99953814, 0.9995667 , 0.99959406,
0.99962037, 0.99964579, 0.99966868, 0.99969089, 0.99971229,
0.99973209, 0.99975141, 0.99976921, 0.99978678, 0.99980322,
0.99981854, 0.99983254, 0.99984585, 0.99985804, 0.99987001,
0.99988166, 0.99989134, 0.99990073, 0.99990983, 0.99991805,
0.99992601, 0.99993322, 0.99994028, 0.99994713, 0.99995323,
0.99995891, 0.99996385, 0.99996848, 0.99997301, 0.99997674,
0.99998028, 0.99998358, 0.99998672, 0.99998936, 0.99999161,
0.99999353, 0.99999532, 0.999997 , 0.9999982 , 0.99999927, 1.

```

3.0.1 Como eu gostaria de manter 80% da variância, o número de dimensões que devo manter deve ser igual a posição do primeiro elemento cujo valor é >0.8. No caso, 13 dimensões.

```
In [8]: pca = sklearn.decomposition.PCA(n_components=13)
```

```
In [9]: new_data_pca = pca.fit_transform(data1)
        print(new_data_pca)
```

```

[[ -2.94087385   9.68333061  -3.62928882 ...,  -1.06968525  -1.13893206
   -0.28967336]
 [  4.7553443  -2.90471677  -4.49673432 ...,   1.6770415  -1.97068645
   -0.57238956]
 [-10.93898353  -3.20602349   2.97069575 ...,   0.0274727   0.23875956
    0.1632103 ]
 ...,
 [ -9.08467378  -3.72857516   0.22389101 ...,  -1.68233633   0.64949083
   -0.69624195]
 [  7.03199293  -3.62465261   1.3818902  ...,  -0.93995011  -0.74491111

```

```

-0.1071369 ]
[  7.17548827 -2.58143506  5.47042863 ..., -0.27656227 -2.58848942
 2.23865287]]

```

3.0.2 Encontramos aqui os novos dados redimensionados pelo PCA

4 2) Logistic Regression

4.0.1 Agora aplicaremos a regressão logística nos dados com PCA

```

In [10]: lr = sklearn.linear_model.LogisticRegression()
         print(lr)

```

```

LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                    verbose=0, warm_start=False)

```

```

In [11]: #Chamamos de X1 e Y1 os dados de treino e teste, respectivamente, com PCA
         X1 =new_data_pca[0:200]
         Y1 = new_data_pca[200:len(new_data_pca)]
         print (X1.shape, Y1.shape)

```

```

(200, 13) (276, 13)

```

```

In [12]: output = (data[1:,-1:])

         treino =(output[0:200, ])
         treino = np.reshape(treino, np.size(treino))

         teste = output[200:len(output),]
         teste = np.reshape(teste, np.size(teste))
         print(treino.shape, teste.shape)

```

```

(200,) (276,)

```

```

In [13]: lr.fit(X1, treino)

```

```

Out[13]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                             penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                             verbose=0, warm_start=False)

```

```

In [14]: pred1 = lr.predict(Y1)

```

```
In [15]: #Calculamos a taxa de acerto da LR com o PCA
         acertol = np.mean(predl == teste)
         print(acertol)
```

0.797101449275

4.0.2 Calcularemos agora a taxa de acerto do LR sem o PCA

```
In [16]: lr2 = sklearn.linear_model.LogisticRegression()
```

```
In [17]: #Chamamos de X2 e Y2 os dados de treino e teste, respectivamente, sem PCA
         X2 = data1[0:200]
         Y2 = data1[200:len(data1)]
         print (X2.shape, Y2.shape)
```

(200, 166) (276, 166)

```
In [18]: #Usamos o conjunto de output já extraído anteriormente
         lr2.fit(X2, treino)
```

```
Out[18]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                             penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                             verbose=0, warm_start=False)
```

```
In [19]: pred2 = lr2.predict(Y2)
```

```
In [20]: #Calculamos a taxa de acerto da LR sem o PCA
         acertol2 = np.mean(pred2 == teste)
         print(acertol2)
```

0.797101449275

4.0.3 Podemos ver que obtemos a mesma taxa de acerto em ambos os métodos

5 3) LDA

5.0.1 Vamos fazer o LDA com os dados do PCA:

```
In [21]: lda1 = lda.LinearDiscriminantAnalysis()
```

```
In [22]: lda1.fit(X1, treino)
```

```
Out[22]: LinearDiscriminantAnalysis(n_components=None, priors=None, shrinkage=None,
                                     solver='svd', store_covariance=False, tol=0.0001)
```

```
In [23]: pred3 = lda1.predict(Y1)
```

```
In [24]: #Calculamos a taxa de acerto do LDA com o PCA
         acerto3 = np.mean(pred3 == teste)
         print(acerto3)
```

0.786231884058

5.0.2 Vamos fazer o LDA sem os dados do PCA

```
In [25]: lda2 = lda.LinearDiscriminantAnalysis()
```

```
In [26]: lda2.fit(X2, treino)
```

```
Out[26]: LinearDiscriminantAnalysis(n_components=None, priors=None, shrinkage=None,
                                     solver='svd', store_covariance=False, tol=0.0001)
```

```
In [27]: pred4 = lda2.predict(Y2)
```

```
In [28]: #Calculamos a taxa de acerto do LDA sem o PCA
         acerto4 = np.mean(pred4 == teste)
         print(acerto4)
```

0.677536231884

6 4) Podemos concluir que o Logistic Regression não possui diferença de desempenho no uso de dados com PCA. No entanto, o LDA tem sua taxa de acerto otimizada.